



PPO fine-tune training trajectories

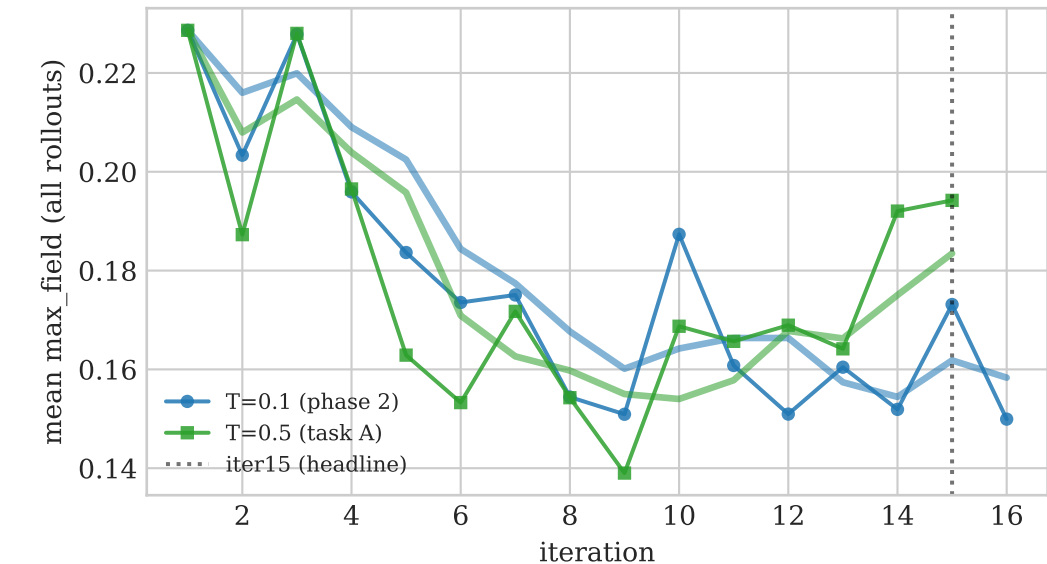
Training return (per iter)



Safety -- p95 max_field per iter (lower is safer)



Safety -- mean max_field per iter



Goal-reaching per iter

